

Mark Rumore

Manager, ML / AI Engineering

markrumore39@gmail.com

(630)450-6230

Chicago, IL, United States

<https://www.linkedin.com/in/mark-rumore/>

<https://markrumore.github.io/>

Technical ML Leader with experience in architecting and managing end-to-end MLOps pipelines, including governance, CI/CD, containerization (Docker, Kubernetes), deployment, and monitoring at scale. Proven track record in developing and productionizing ML, deep learning, and NLP models on Azure and AWS. Skilled in leading cross-functional teams, mentoring engineers, and driving Agile practices. Experienced in building real-time ml pipelines and streaming architectures for low-latency AI applications. Expertise in AI strategy, model compression, active learning, knowledge distillation, distributed training, feature stores, automated retraining, transfer learning, and interpretability to deliver robust, scalable AI solutions.

Work Experience

Manager, AI Engineering

Nov 2025 – Present

PwC – Chicago, IL

Technical Lead, Senior ML Engineer

May 2023 – Nov 2025

Coca-Cola Bottler – Chicago, IL

- Built departmental MLOps pipelines, implementing governance best practices, model lifecycle management, automated CI/CD workflows, and code review processes (Azure Pipelines, Docker, API Endpoints) to cut development and scalability time by 50% and reduce code vulnerabilities by 30%.
- Chaired the AI Governance Review Board, driving enterprise policies for responsible AI and ensuring 100% compliance with ethical and industry standards, including bias mitigation, fairness assurance, and model transparency frameworks.
- Directed cross-functional teams to implement scalable, production-grade ML systems on cloud-native architectures (Azure ML, Databricks, Snowflake, Kafka); launched the company's largest production model, reducing employee labor hours by 30% through outlet-level process optimization.
- Architected and deployed enterprise-scale models (CNN, ViT, RNN, real-time systems, APIs), transformer architectures, and recommendation systems, delivering solutions with optimal performance and model interpretability.

Data Scientist

Feb 2021 – May 2023

XSELL Technologies – Chicago, IL

- Engineered NLP models for chat/voice cobots, enhancing speech pattern recognition to optimize call agent responses.
- Led AT&T client model development, applying tuning and calibration techniques to prevent overfitting and ensure robust performance. (Deep Learning, fine tuning, Python, AWS, S3)
- Applied predictive modeling to refine classifiers, improving voice data capture efficiency in workflow processes. (refactor)
- Optimized ML Flow pipelines, including multinomial classifier calibration, transformer models, active learning, and few-shot learning techniques. (Transformers, hugging face, SageMaker)

Data Analyst

Feb 2019 - Feb 2021

Allied Benefit Systems – Chicago, IL

- Developed a departmental Access database to centralize industry-specific data, enabling comprehensive risk analysis and strategic development. (SQL, DBM, Actuarial Methods)

Core Skills

Programming: Python, R, SQL, Java, C - **Cloud/MLOps:** AWS, Azure, Databricks, Snowflake - **Container:** Docker, Kubernetes

CI/CD: Jenkins, GitHub **Workflow:** MLflow, Kubeflow, Terraform - **Additional Skills:** Agile/Scrum, Real-time deployments, ETL

Visualization: PowerBi, Tableau, DAX, Sisense

Education

University of Colorado Boulder – Master of Science - Computer Science

DePaul University – Master of Science - Data Science

2022

University of Iowa – Bachelor of Arts - Economics

2018